Hands-on “Vectorization”

- Automatic Vectorization
  - Vectorization Report
- Guided Vectorization
  - pragma simd
  - pragma ivdep
  - restrict
- Intel Advisor
  - Create Advisor Project
  - Collect Survey Data
  - Change xhost
- Matrix Multiplication
  - New advisor Project
  - Collect Survey Data
  - Check Trip Count
  - Check Dependencies
  - Memory Access Patterns

1) Compile the example with vec-report6 and O3

```
icc VectorizationHandson.c -o VectorizationHandson -vec-report6 –O3
```

2) Open the vectorization report (VectorizationHandson.optrpt)

Note that the loop on function main was automatically vectorized;

remark #15300: LOOP WAS VECTORIZED

Note that the loop on function add_floats was not automatically vectorized;

remark #15344: loop was not vectorized: vector dependence prevents vectorization

2) Include “#pragma ivdep directive” in top of loop “for (i=0; i<n; i++)“ on function add_floats

```
#pragma ivdep
for (i=0; i<n; i++){
```

3) Recompile the example with vec-report

```
icc VectorizationHandson.c -o VectorizationHandson -vec-report6 –O3
```

Note that the Outer loop can be vectorized now.

remark #15301: OUTER LOOP WAS VECTORIZED

2) Include pragma simd directive in the loop “for (j=n; j>0; j--){“ on function add_floats

```
#pragma simd
for (j=n; j>0; j--){
```

3) Recompile the example with vec-report

icc VectorizationHandson.c -o VectorizationHandson -vec-report6 –O3

Note that compiler vectorized the inner loop instead of outer loop

Outer loop:
remark #15542: loop was not vectorized: inner loop was already vectorized
Inner loop:
remark #15301: SIMD LOOP WAS VECTORIZED

2) Include keyword restrict to avoid in all arguments of function quad

void quad(int length, double * restrict a, double * restrict b, double * restrict c, double * restrict x1, double * restrict x2)

3) Recompile the example with vec-report and restrict

icc VectorizationHandson.c -o VectorizationHandson -vec-report6 –O3 -restrict

Note that the loop can be vectorized now.

remark #15300: LOOP WAS VECTORIZED

Intel Advisor – Create New Project

- Execute Intel Advisor on terminal: advixe-gui
- create new Advisor project:
  - name: Vectest
  - application: ~/handson/vec/VectorizationHandson
  - Source Folder: ~/handson/vec/

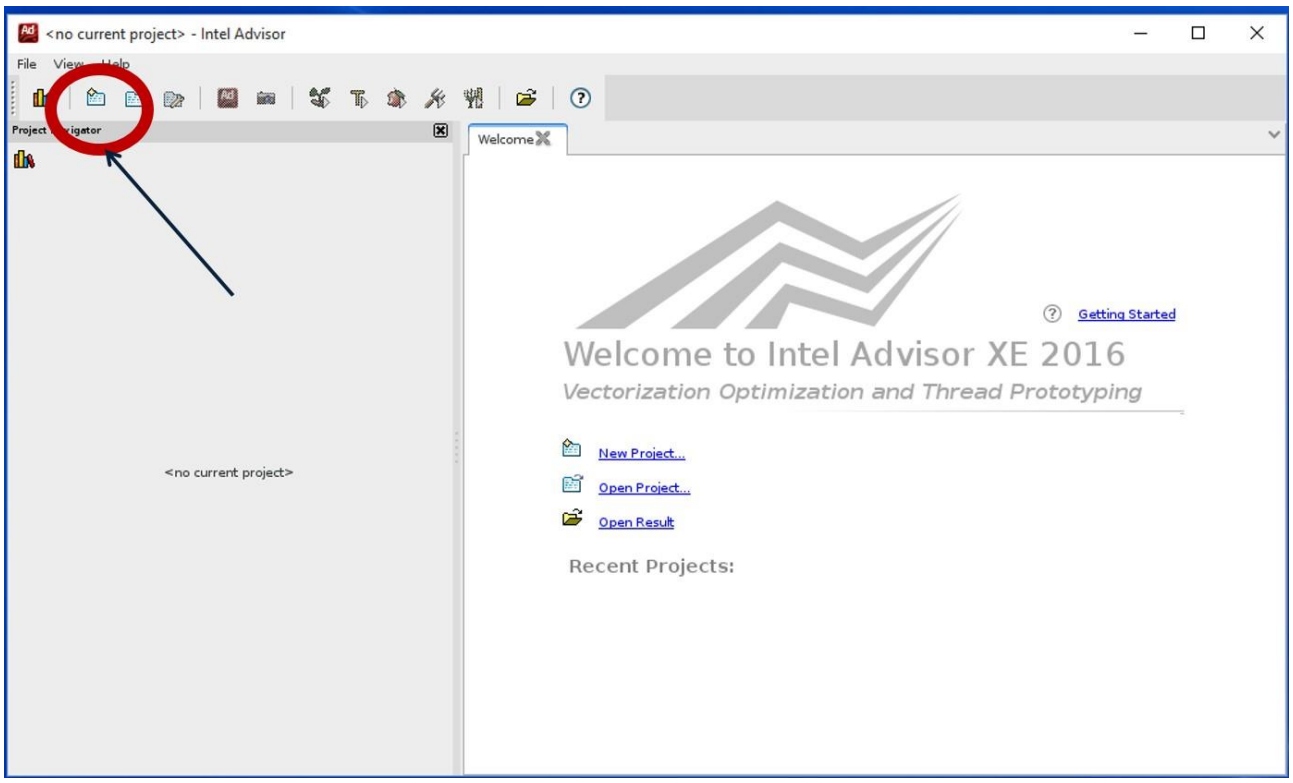
<no current project> - Intel Advisor
File View Help
Project Navigator
<no current project>
Welcome
Getting Started
Welcome to Intel Advisor XE 2016
Vectorization Optimization and Thread Prototyping
New Project...
Open Project...
Open Result
Recent Projects:

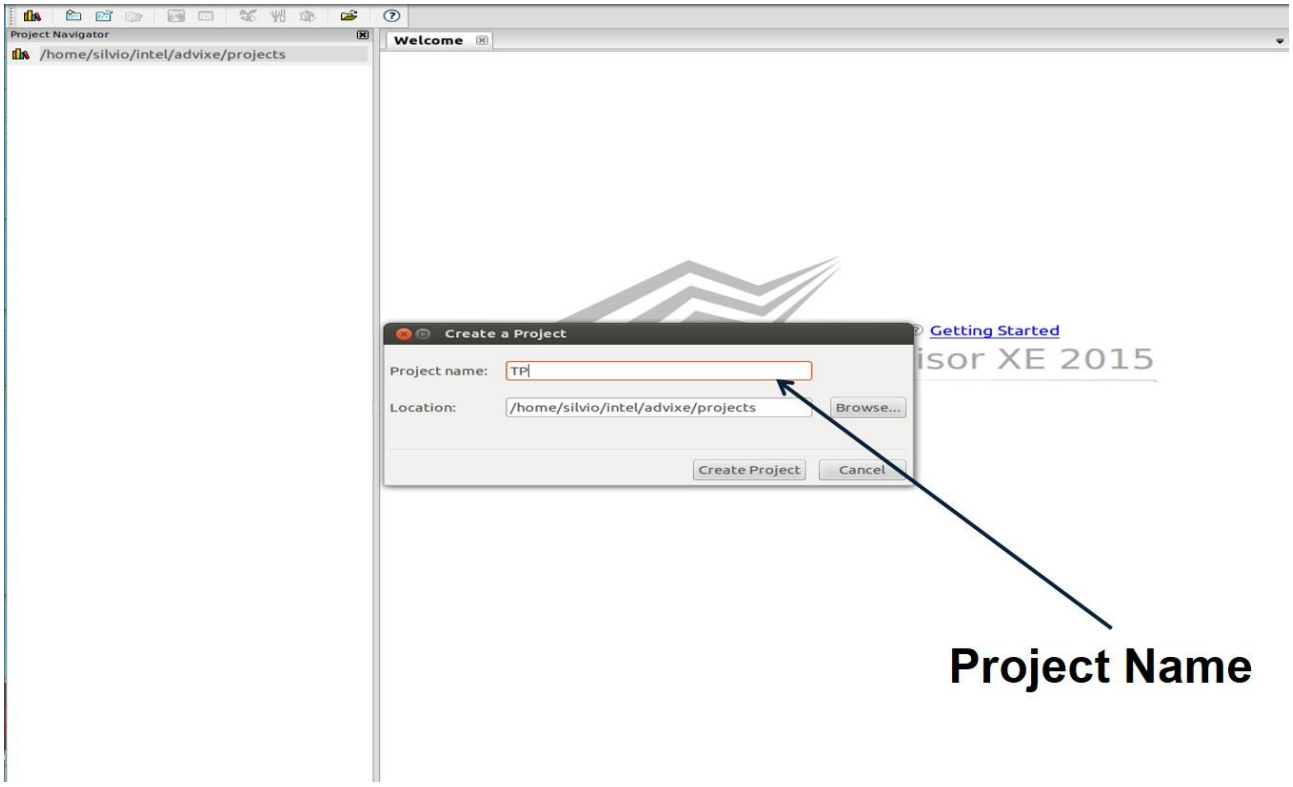
Project Navigator
/home/silvio/intel/advixe/projects
Welcome
Create a Project
Project name: TP
Location: /home/silvio/intel/advixe/projects
Browse...
Create Project
Cancel
Getting Started
isor XE 2015
Project Name

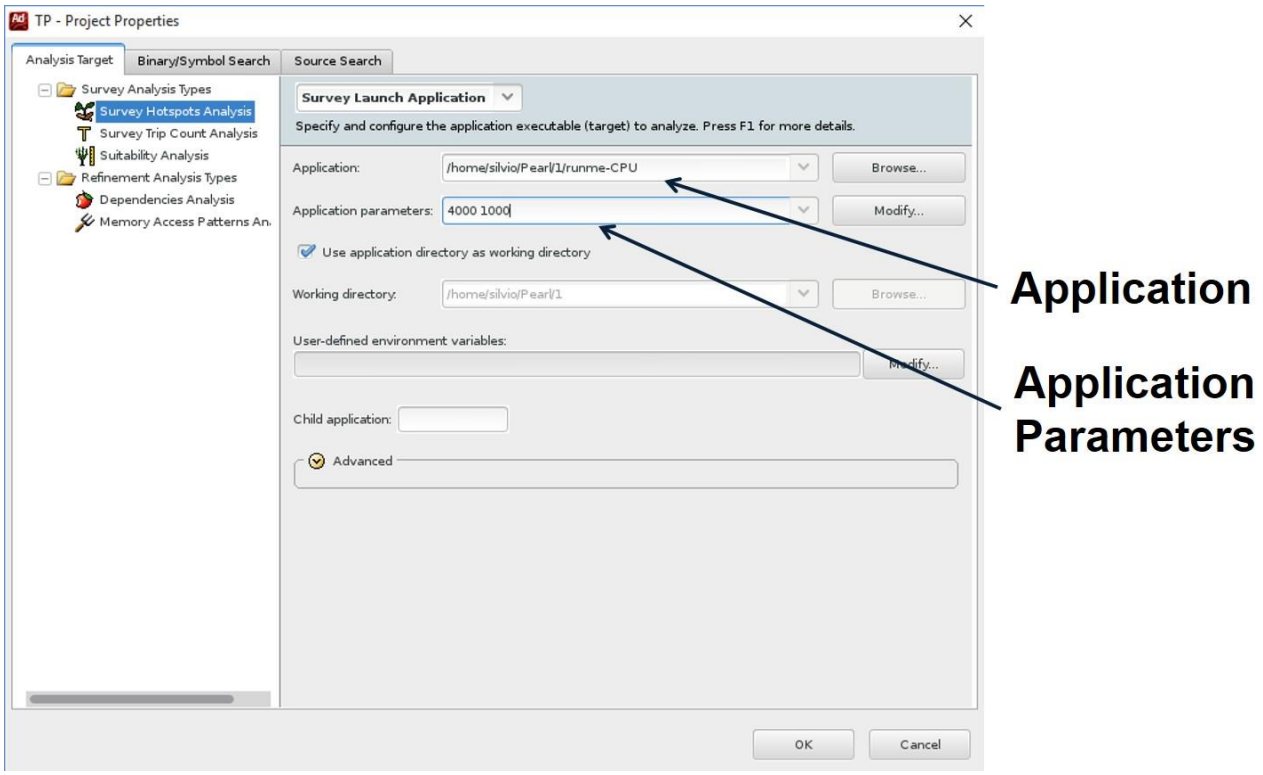
TP - Project Properties
Analysis Target
Binary/Symbol Search
Source Search
Survey Analysis Types
Survey Hotspots Analysis
Survey Trip Count Analysis
Suitability Analysis
Refinement Analysis Types
Dependencies Analysis
Memory Access Patterns An.
Survey Launch Application
Specify and configure the application executable (target) to analyze. Press F1 for more details.
Application:
/home/silvio/Pearl/1/runme-CPU
Browse...
Application parameters:
4000 1000
Modify...
Use application directory as working directory
Working directory:
/home/silvio/Pearl/1
Browse...
User-defined environment variables:
Modify...
Child application:
Advanced
OK
Cancel
Application
Application Parameters

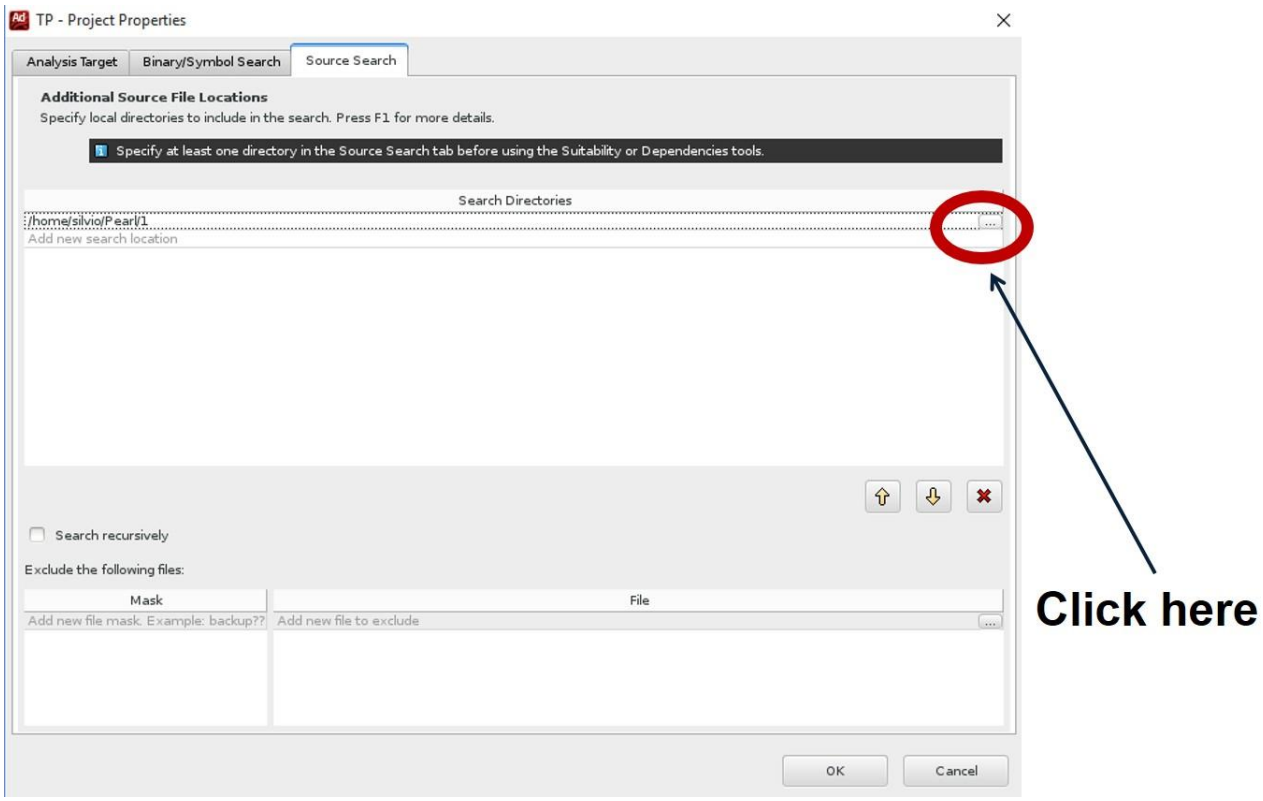
TP - Project Properties
Analysis Target
Binary/Symbol Search
Source Search
Additional Source File Locations
Specify local directories to include in the search. Press F1 for more details.
Specify at least one directory in the Source Search tab before using the Suitability or Dependencies tools.
Search Directories
/home/silvio/Pearl/1
Add new search location
Search recursively
Exclude the following files:
Mask
File
Add new file mask. Example: backup??
Add new file to exclude
OK
Cancel
Click here

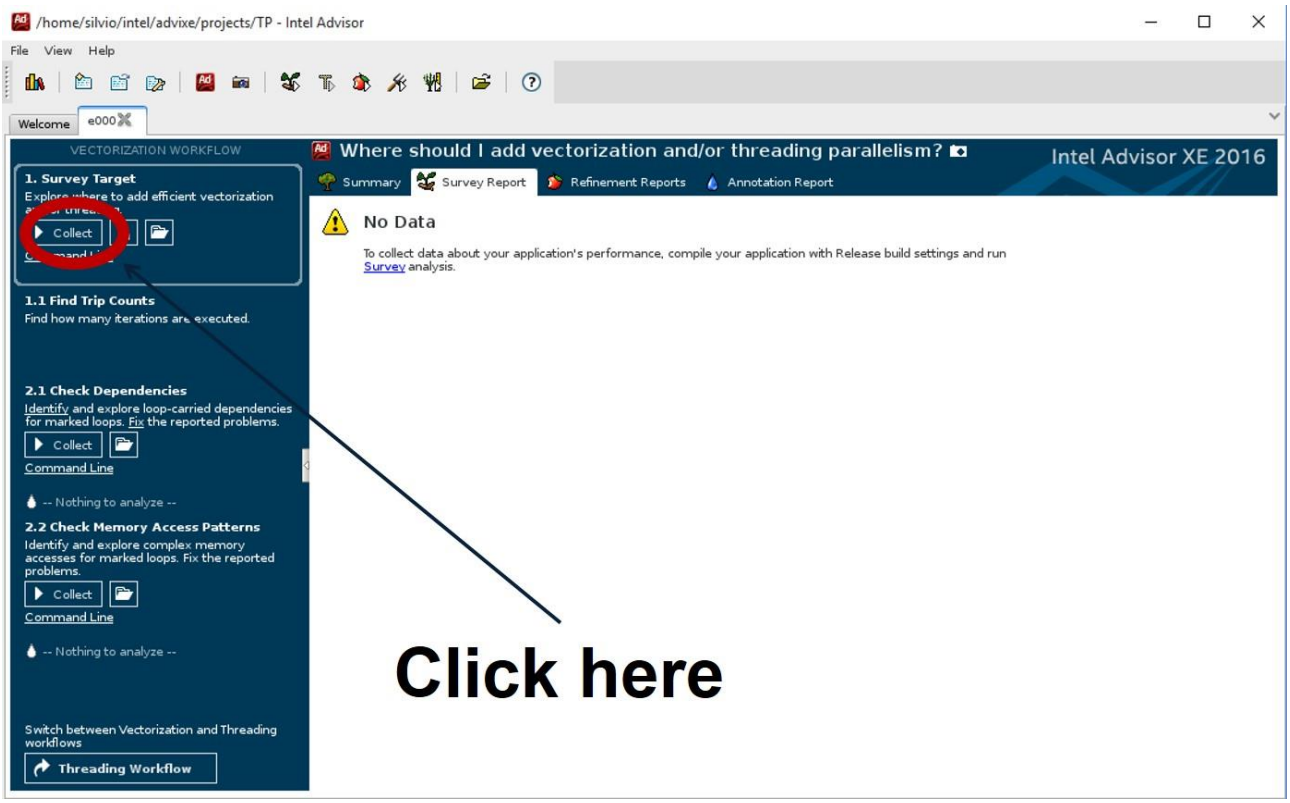

Change xhost

- Note that the code was compiled using SSE
- Recompile application using xhost and –g

icc VectorizationHandson.c -o VectorizationHandson -vec-report6 –O3 –restrict –xhost -g

- Collect Survey Data again
- Note that now the code was compiled using AVX

Matrix Multiplication

- Compile code;
- Create new Advisor Project;
- Execute Survey Analysis;
- Execute Trip Count Analysis;
- Execute Check Dependencies Analysis;

- Compile code
  - Cd ~/handson/matrix/linux
  - Make clean ; make icc
- Execute Intel Advisor on terminal: advixe-gui
- create new Advisor project:
  - name: matrix-handson
  - application: ~/handson/matrix/linux/matrix.icc
  - Source Folder: ~/handson/matrix/src/
- Execute Survey Analysis

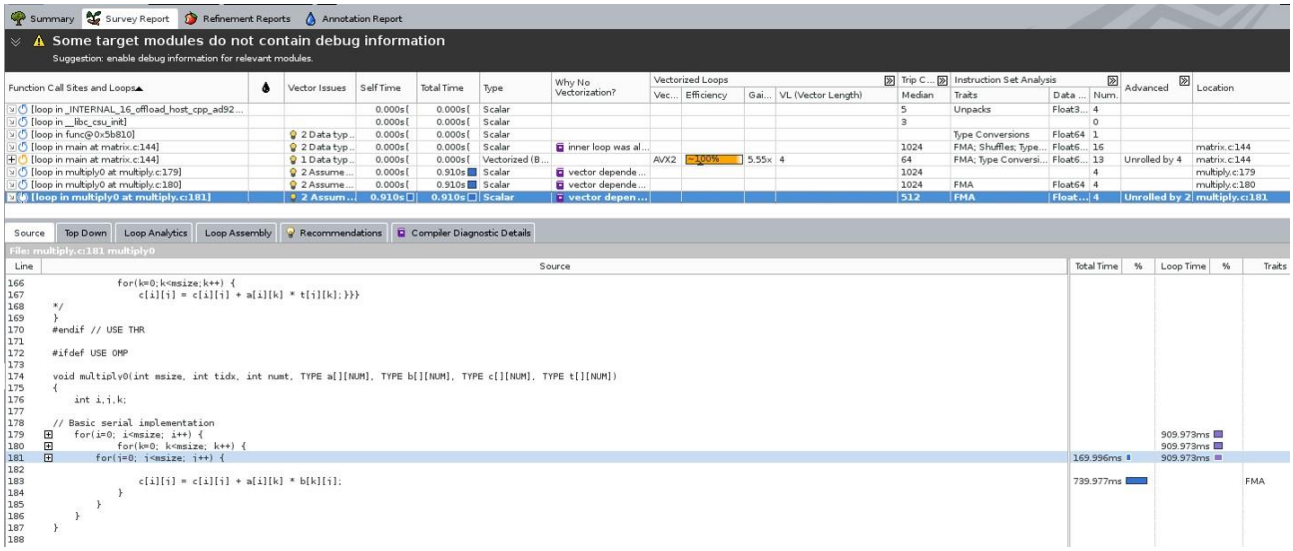

- Note that Inner loop was not automatic vectorized.

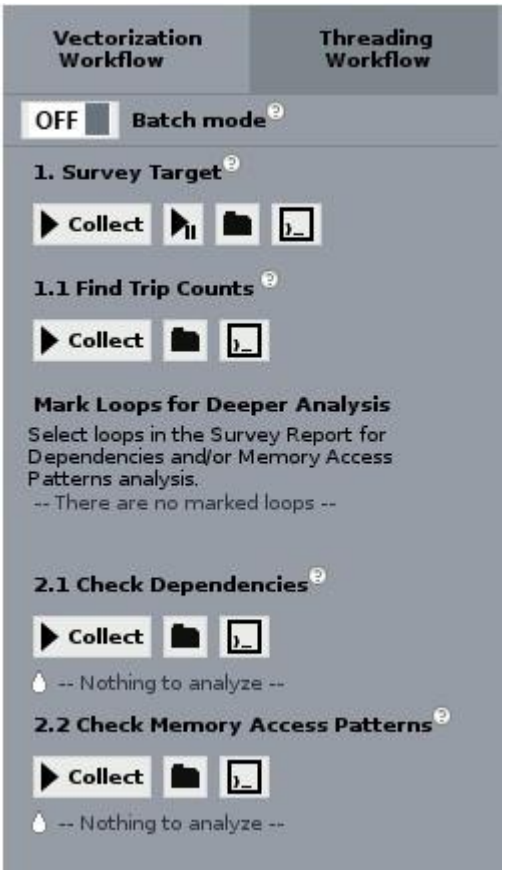

Trip Counts

Median	Min	Max	Call Count	Iteration Duration
5	5	5	1	
3	3	3	1	
1024	1024	1024	1	
64	64	64	1024	
1024	1024	1024	1	0.0009s
1024	1024	1024	1024	< 0.0001s
512	**512**	**512**	**1048576**	**< 0.0001s**

- Mark Inner Loop for deeper analysis;
- Click on “check dependency”;

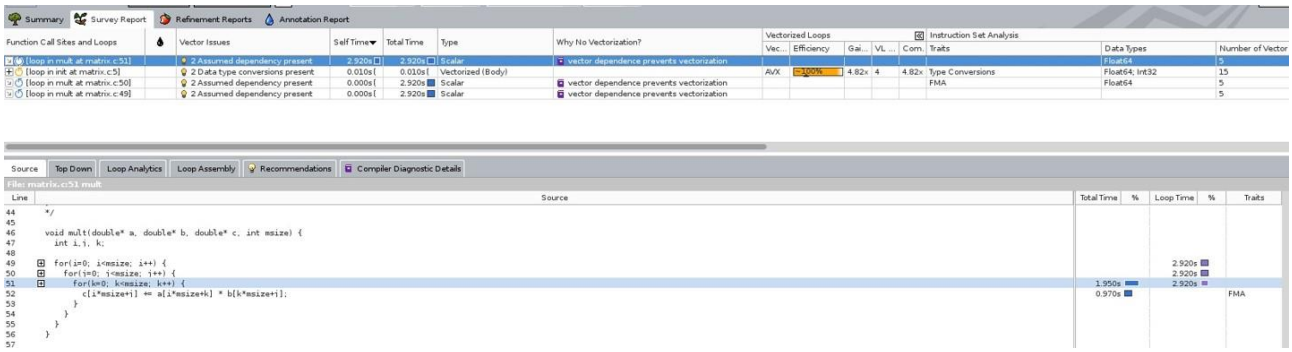

- No dependencies found. It is safe to vectorize;

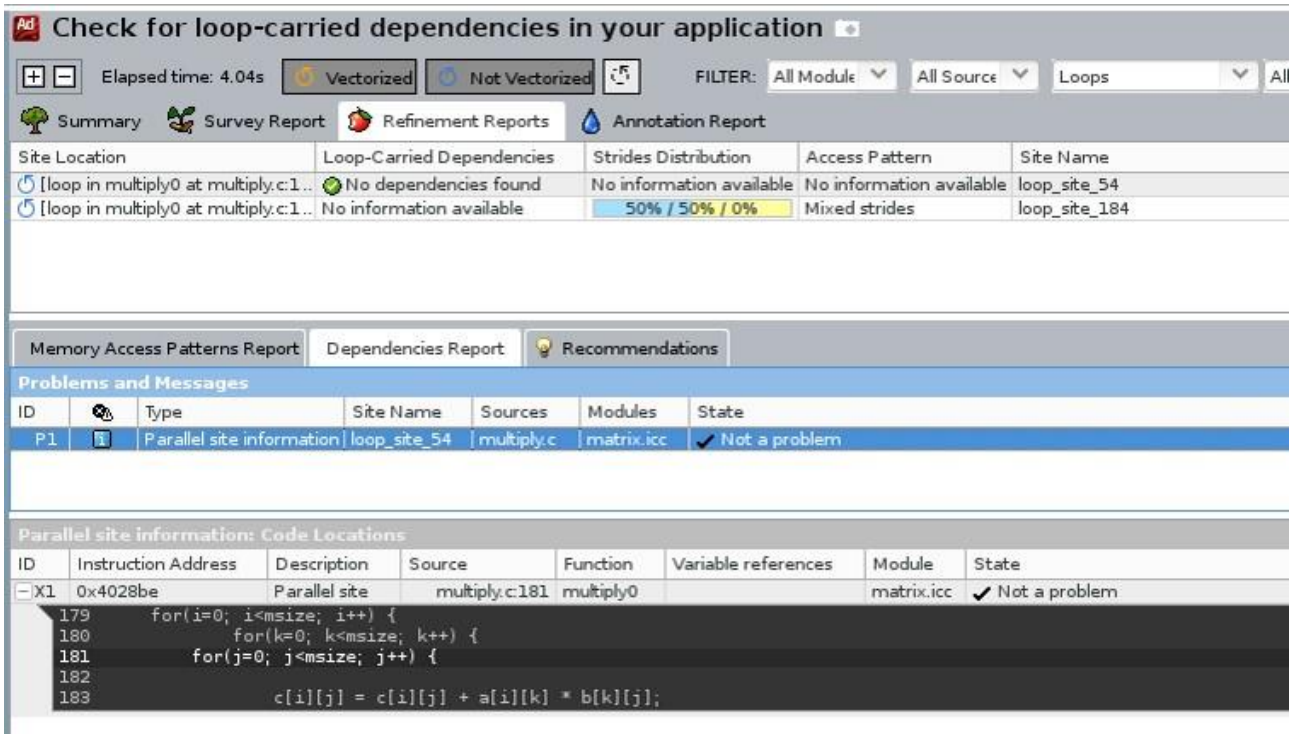

- Put #pragma simd in top of inner loop on function multiply0:
  - Nano ~/handson/matrix/src/multiply.c

```
for(i=0; i<msize; i++) {
   for(k=0; k<msize; k++) {
      #pragma simd
  for(j=0; j<msize; j++) {
```

- Recompile application
  - Make clean ; make icc
  - Run survey data again;
- Note that inner loop is vectorized now;

- Run Check Memory Access Patterns;

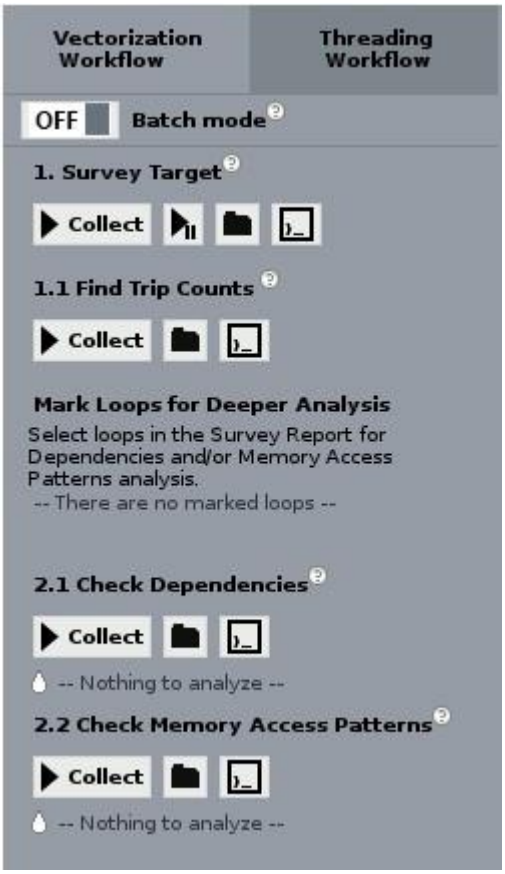

- Memory access Patterns Results:
  - 50% constant stride;
  - 50% unit stride;

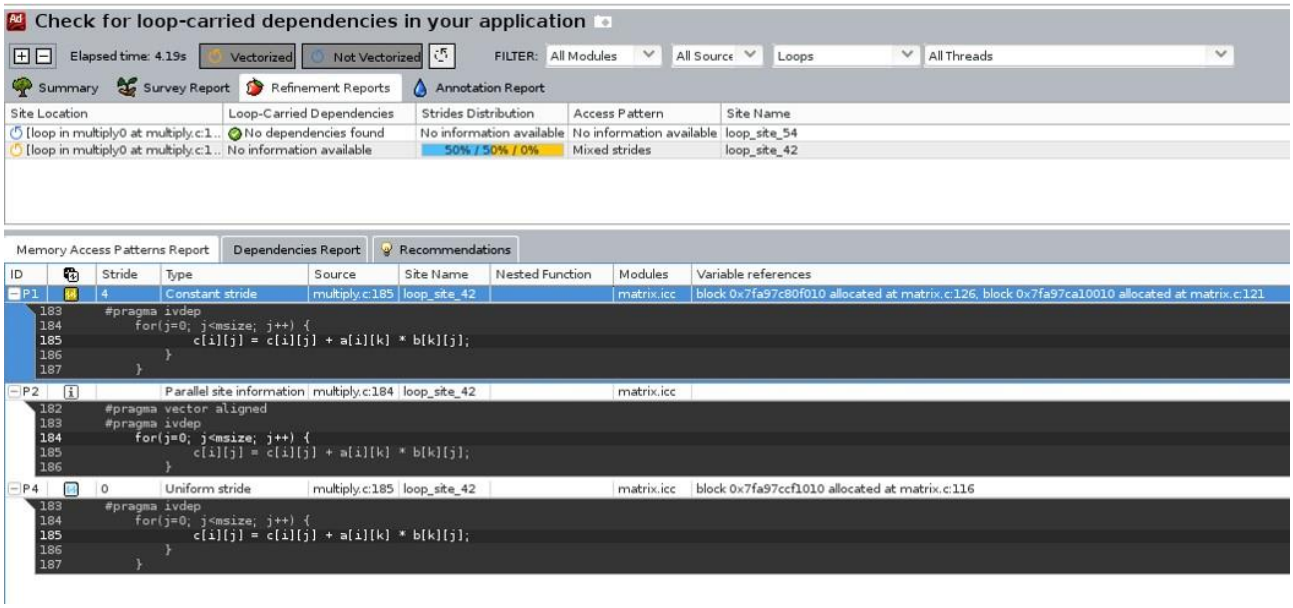
Check for loop-carried dependencies in your application
Elapsed time: 4.19s
Vectorized
Not Vectorized
FILTER:
All Modules
All Source
Loops
All Threads
Summary
Survey Report
Refinement Reports
Annotation Report
Site Location
Loop-Carried Dependencies
Strides Distribution
Access Pattern
Site Name
[loop in multiply0 at multiply.c:1...
No dependencies found
No information available
No information available
loop_site_54
[loop in multiply0 at multiply.c:1...
No information available
50% / 50% / 0%
Mixed strides
loop_site_42
Memory Access Patterns Report
Dependencies Report
Recommendations
ID
Stride
Type
Source
Site Name
Nested Function
Modules
Variable references
P1
4
Constant stride
multiply.c:185
loop_site_42
matrix.icc
block 0x7fa97c80f010 allocated at matrix.c:126, block 0x7fa97ca10010 allocated at matrix.c:121
183 #pragma ivdep
184 for(j=0; j<msize; j++) {
185 c[i][j] = c[i][j] + a[i][k] * b[k][j];
186 }
187 }
P2
Parallel site information
multiply.c:184
loop_site_42
matrix.icc
182 #pragma vector aligned
183 #pragma ivdep
184 for(j=0; j<msize; j++) {
185 c[i][j] = c[i][j] + a[i][k] * b[k][j];
186 }
P4
0
Uniform stride
multiply.c:185
loop_site_42
matrix.icc
block 0x7fa97ccf1010 allocated at matrix.c:116
183 #pragma ivdep
184 for(j=0; j<msize; j++) {
185 c[i][j] = c[i][j] + a[i][k] * b[k][j];
186 }
187 }